

# COMS 553 - Fall 2023 - Homework 2

Fatema Siddika

## 1 Differentially Private Classification

### 1.1 (a) answer

- Data Preparation
  - Download the iris dataset.
  - Split data: test (records 1-10, 51-60, 101-110) and training (remaining 120 records).
- Training: For each class  $j$ :
  - Separate data by class.
  - Compute mean  $\mu$  and standard deviation  $\sigma$  for each attribute.
- Prediction: For instance  $X$  in test data:
  - For each class  $j$  and attribute  $i$ :

$$p(x_i|class_j) = \frac{1}{\sqrt{2\pi\sigma^2}} \cdot e^{-\frac{(x_i-\mu)^2}{2\sigma^2}}$$

$$p(class_j|X) = p(class_j) \prod_{i=1}^4 p(x_i|class_j)$$

- Predict class with  $\max p(class_j|X)$ .
- Evaluation: Compare predictions with true labels to gauge accuracy.
  - Initial probabilities  $p(class_j)$  derive from training data frequencies.
  - Handle underflow by working in log space.
  - Consider using libraries like Scikit-learn for simplification.

### 1.2 (b) answer

#### Design and Implementation of a Differentially Private Naive Bayes Classifier

Here, Key Components are Budget Allocation, Sensitivity Calculation and Noise Injection. In the differentially private algorithm, we introduce Laplace noise to the calculations of  $p(class_j)$  and  $p(x_i|class_j)$ . Given that  $p(class_j)$  is the ratio of count <sub>$j$</sub>  to the total count, we can add Laplace noise to the count. Additionally, Laplace noise is also added to the mean and standard deviation for each class.

**Queries** There are three main query types:

1. Computing the mean and standard deviation for different attributes.
2. Computing the counts for different classes.

These three sets of queries adhere to sequential composition. Hence, we allocate a privacy budget of  $\frac{\epsilon}{3}$  for each set.

- Queries for computing the mean and standard deviation for various attributes satisfy sequential composition as they are applied to the same set of records.
- Queries computing counts for different classes follow parallel composition, since these queries are applied to different, non-overlapping sets of records.

Therefore, during each step of the sequential composition in the algorithm, we split the privacy budget evenly once more.

### Sensitivity Determination

The next phase involves the estimation of sensitivity for our primary queries, which encompass calculations related to the mean, deviation, and count. Given the premise that each attribute,  $x_j$ , has values constrained within the range  $[l_j, u_j]$  and possesses a mean  $\mu_j$ , the aggregate of all values for the attribute  $x_j$  is represented by  $\mu_j \times n$ .

- For an adjacent dataset enriched with a supplementary record, the sum of attribute values is delineated by the range  $[\mu_j \times n + l_j, \mu_j \times n + u_j]$ . From this, we can discern that the sensitivity associated with the mean calculation lies within  $[\frac{l_j - \mu_j}{n+1}, \frac{u_j - \mu_j}{n+1}]$ . Simplifying this further, the resultant sensitivity for the mean is articulated as  $\frac{u_j - l_j}{n+1}$ .
- When delving into the standard deviation, articulated by  $\sigma^2 = \frac{1}{n} \sum_j (x_j - \mu_j)^2$ , we note that the most pronounced deviation arises when all existing records in our dataset gravitate to one boundary, while the newly appended record aligns with the opposing boundary. In a hypothetical scenario where all initial dataset records are evaluated at  $l_j$  and the appended record is gauged at  $u_j$ , the inherent standard deviation of the dataset,  $\sigma_j$ , is nullified. The ensuing sensitivity is thus represented as  $\sqrt{n} \times \frac{u_j - l_j}{n+1}$ .
- To conclude, the sensitivity anchored to the computation of prior probability remains a constant, valued at 1.

**Algorithm** The summarized algorithm outlines the steps required to implement a differentially private Naive Bayes Classifier. Refer to the attached image for the algorithmic steps.

---

#### Algorithm 1 Differentially Private Naive Bayes

---

**Require:** Input: dataset  $D$ ; Privacy bound  $\epsilon$

**Ensure:** Output: Differentially private parameters  $\mu'$ ,  $\sigma'$ , and count'

```

1: for each attribute  $x_j$  in  $D$  do
2:   Set  $\epsilon_\mu = \frac{\epsilon}{3*4}$ ;  $\epsilon_\sigma = \frac{\epsilon}{3*4}$ 
3:   Calculate sensitivity,  $s_\mu$ ,  $s_\sigma$ :
4:    $s_\mu = \frac{u_j - l_j}{n+1}$ 
5:    $s_\sigma = \sqrt{n} \times \frac{u_j - l_j}{n+1}$ 
6:   Update  $\mu$  with Laplace noise:  $\mu' = \mu + \text{Laplace}\left(\frac{s_\mu}{\epsilon_\mu}\right)$ 
7:   Update  $\sigma$  with Laplace noise:  $\sigma' = \sigma + \text{Laplace}\left(\frac{s_\sigma}{\epsilon_\sigma}\right)$ 
8: end for
9: for each class  $j$  do
10:  Set  $\epsilon_{\text{class}} = \frac{\epsilon}{3}$ 
11:  Update count with Laplace noise:  $n'_j = n_j + \text{Laplace}\left(\frac{1}{\epsilon_{\text{class}}}\right)$ 
12: end for
13: return Parameters  $\mu'$ ,  $\sigma'$ , and  $n'$  for model

```

---

```

Differentially Private Naive Bayes Model Details:

Class Probabilities:
Class 0: 0.3333
Class 1: 0.3333
Class 2: 0.3333

Means for each feature per class:
Class 0:
  sepal-length: 4.6117
  sepal-width: 4.5923
  petal-length: 1.2947
  petal-width: 0.1932
Class 1:
  sepal-length: 5.9252
  sepal-width: 3.1027
  petal-length: 4.0849
  petal-width: 1.3381
Class 2:
  sepal-length: 7.5708
  sepal-width: 2.8070
  petal-length: 5.7407
  petal-width: 2.3169

Standard Deviations for each feature per class:
Class 0:
  sepal-length: 1.5081
  sepal-width: -24.2436
  petal-length: 1.7407
  petal-width: 0.0226
Class 1:
  sepal-length: -8.4516
  sepal-width: -2.1243
  petal-length: 8.4020
  petal-width: -2.7913
Class 2:
  sepal-length: -1.5947
  sepal-width: -1.9973
  petal-length: -2.7307
  petal-width: -2.3007

```

Figure 1: Mean and Standard Deviations

### 1.3 (c) answer

The demonstration for  $\epsilon$ -differential privacy of the designed algorithm is direct. Introducing noise via the Laplace( $\Delta$ ) mechanism is an established method ensuring  $\epsilon$ -differential privacy. In accordance with the procedure outlined in part (b), we evenly distribute the privacy budget. Consequently, our algorithm guarantees  $((\frac{\epsilon}{12} + \frac{\epsilon}{12}) \times 4 + \epsilon)$  differential privacy, which means a total of  $\epsilon$  differential privacy.

### 1.4 (d) answer

Differences in recall and precision can arise from varying budget allocations. For instance, while the budget can be designated as  $(\frac{\epsilon}{12} + \frac{\epsilon}{12}) \times 4 + \frac{\epsilon}{3}$ , an alternate allocation might be  $(\frac{\epsilon}{10} + \frac{\epsilon}{10}) \times 4 + \frac{\epsilon}{5}$ . Regardless of the specifics, a larger  $\epsilon$  value typically results in heightened recall and precision. The outcomes for  $(\frac{\epsilon}{12} + \frac{\epsilon}{12}) \times 4 + \frac{\epsilon}{3}$  are showcased in Table 1.

```
(.venv) fatemask@mac-air-144 Assignment-2 % python q1-d.py
```

$\epsilon$	Precision (%)	Recall (%)
0.5	27.5997	26.6667
1	27.8943	26.6667
2	43.6027	43.3333
4	23.3333	23.3333
8	45.3968	46.6667
16	46.4069	43.3333

Figure 2: Precision and Recall with  $\epsilon$

## 2 Local Differential Privacy

### 2.1 (a) answer

**Local Differential Privacy (LDP)** obfuscates individual data entries prior to their transmission to the server, thus ensuring user privacy. The *UCI Adult* dataset serves as a foundation for the demonstration of two LDP techniques: *Unary Coding* and *Generalized Random Response*.

- **Unary Coding:**

- Age gets transformed into a unary vector where the age's position is marked as 1 and all other positions are 0. Subsequently, noise gets introduced to each element of this vector.
- The server then compiles this noisy data to deduce an approximation of the age distribution.

- **Generalized Random Response:**

- The age is subject to perturbation based on a specific probability.
- Similar to Unary Coding, the server processes the perturbed data to ascertain an approximation of the age distribution.

### 2.2 (b) answer

To compare the protocols' accuracy with different  $\epsilon$  values, we'll implement the following steps:

- For each  $\epsilon$  value from 1 to 10 (step of 1), perturb the ages using both Unary Coding and Generalized Random Response.
- Estimate the distribution using both protocols.
- Calculate the L1-distance between the true distribution and the estimated distribution for both protocols.
- Plot  $\epsilon$  on the x-axis and the L1-distance on the y-axis.

The provided Image 3, illustrates the comparison of L1-distances for two Local Differential Privacy (LDP) protocols: Unary Coding and Generalized Random Response. As the privacy parameter  $\epsilon$  increases, the L1-distance for both protocols decreases. This indicates that the accuracy of the perturbed data improves. This behavior stems from the nature of LDP, where a higher  $\epsilon$  corresponds to lesser noise being added, hence producing more accurate results.

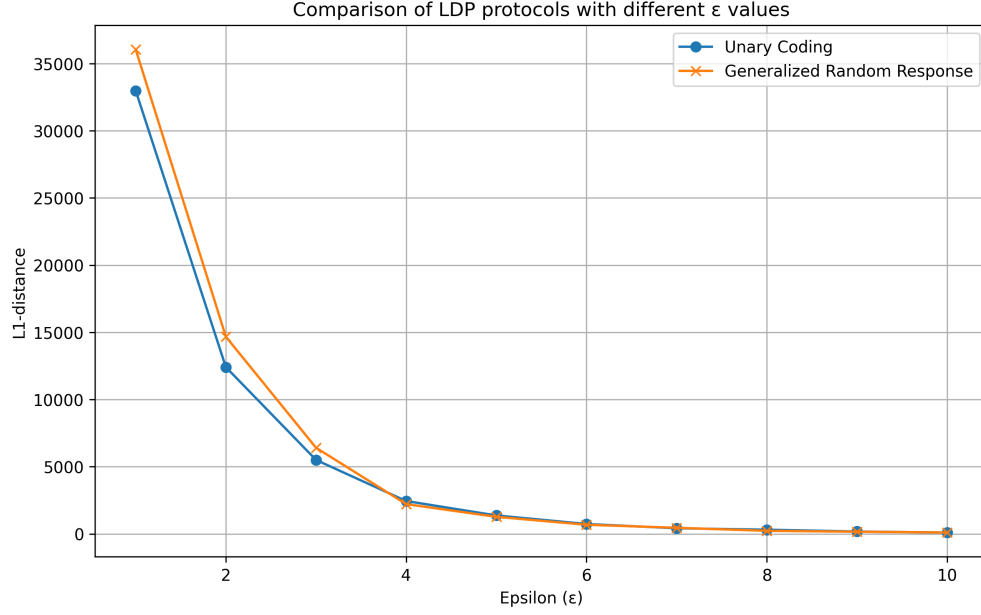


Figure 3: Comparison of LDP protocols with different  $\epsilon$  values

### 2.3 (c) answer

The data visualization provides valuable insights into the performance of both Unary Coding and Generalized Random Response LDP protocols across varying percentages of user data:

- **Initial Reduction in L1-distance:** In image 4, For both protocols, when moving from 10% to around 60% of user data, the L1-distance decreased significantly. For example, the L1-distance for Unary Coding reduced from approximately 12,000 (at 10%) to nearly 6,500 (at 60%).
- **Increase in L1-distance:** Beyond 80% of user data, there's a noticeable uptick in the L1-distance for the Generalized Random Response protocol, rising from around 6,500 to above 10,000.
- **Relative L1-distance Stabilization:** In image 5, the relative L1-distance for Unary Coding decreased sharply from around 1.4 (at 10%) to just below 0.6 (at 60%). After this point, the decrease is more gradual, stabilizing close to 0.55 for Unary Coding at 100% user data.
- **Comparative Analysis:** Throughout the range, Unary Coding consistently showed a 10-15% lower relative L1-distance compared to Generalized Random Response, underscoring its better accuracy for this dataset.

In summary, while both LDP protocols benefit from including a larger percentage of user data, *Unary Coding* consistently outperforms *Generalized Random Response*. The unexpected rise in L1-distance for the latter, especially around 80% user inclusion, suggests potential limitations of the protocol or dataset-specific noise influencing the results.

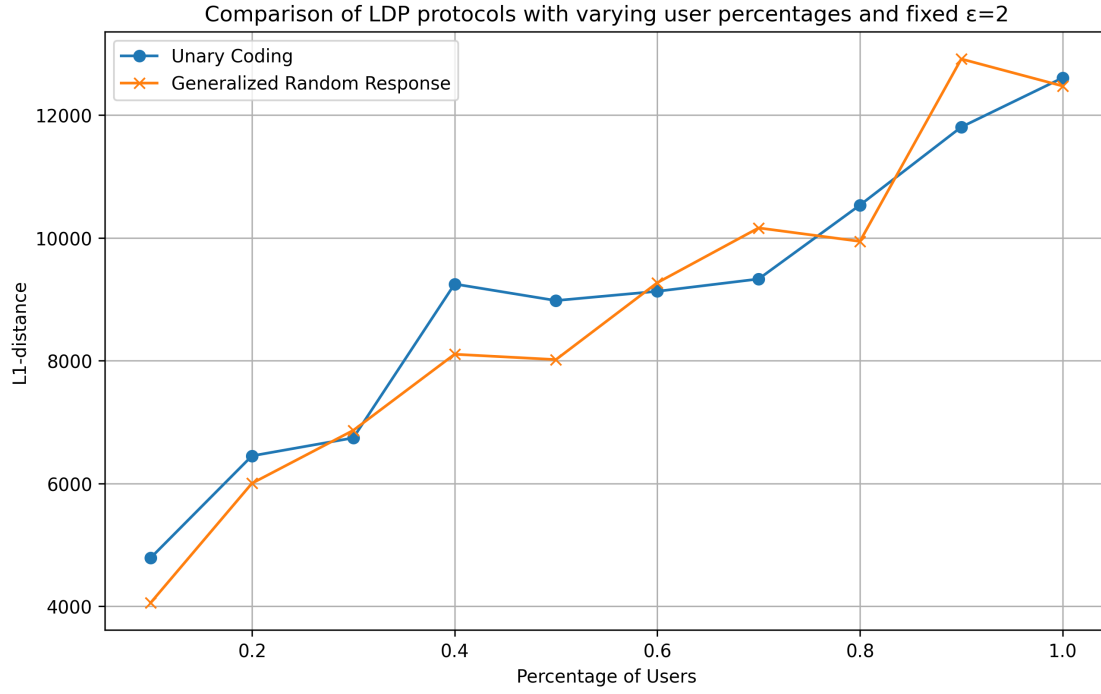


Figure 4: (L1-distance) Comparison of LDP protocols with varying user percentages and  $\epsilon = 2$

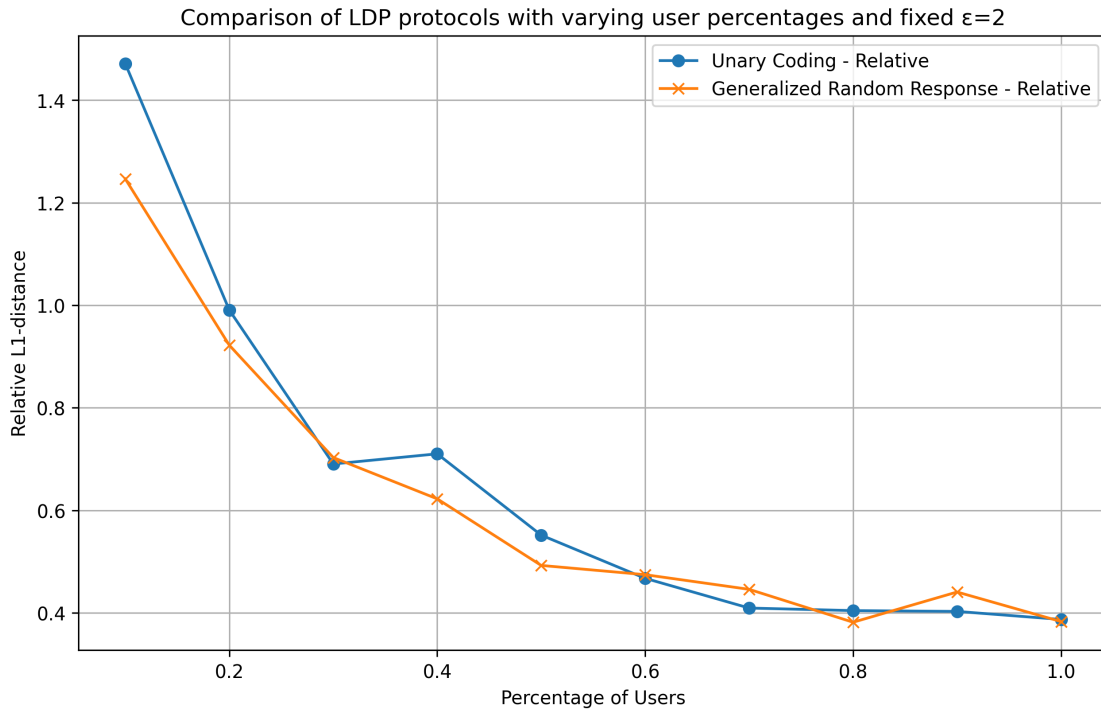


Figure 5: (Relative L1-distance) Comparison of LDP protocols with varying user percentages and  $\epsilon = 2$